

# Shubham Kaushik

[shubhamkaushik.com](https://shubhamkaushik.com) | [LinkedIn](#) | [GitHub](#) | [kaushiks\[at\]brandeis\[dot\]edu](mailto:kaushiks[at]brandeis[dot]edu)

## EDUCATION

Jan 2024 - Present	<b>Doctor of Philosophy</b> in <i>Computer Science</i> (GPA: 3.97/4.0) <a href="#">Brandeis University</a> , MA, United States
Sep 2022 - Dec 2023	<b>Master of Science</b> in <i>Computer Science “Data-Centric Computing”</i> (GPA: 3.88/4.0) <a href="#">Boston University</a> , MA, United States
Jul 2014 - Jun 2018	<b>Bachelor of Technology</b> in <i>Computer Science &amp; Engineering</i> <a href="#">Maharshi Dayanand University</a> , India

## EXPERIENCE

Jan 2024 - Present	<b>Ph.D. Researcher</b> @ Smart & Scalable Data Systems Lab, <a href="#">Brandeis University</a>
Mar 2022 - Aug 2022	<b>Software Engineer</b> @ Server Programming Team, <a href="#">Kwalee</a> , India
Jun 2021 - Mar 2022	<b>Engineer - Information Security</b> @ Cyber Fusion, <a href="#">FIS Global</a> , India
Jul 2018 - Jun 2021	<b>Project Engineer</b> @ Cloud Computing - Wipro Digital, <a href="#">Wipro Limited</a> , India

## TECHNICAL SKILLS

- **Programming Languages:** C, C++, Python, SQL, Rust
- **Markup Languages:** HTML, CSS, JSON, YAML,  $\LaTeX$ , Markdown
- **Databases:** RocksDB, Postgres, MongoDB, Redis, SQLite, ORM
- **Tools & Systems:** Django, Django-REST, Kafka, Hadoop, gRPC, Asyncio, Git, ETL, Flink, AWS

## KEY PROJECTS

- Tectonic** [[Paper PDF](#)] [[GitHub](#)] **Apr 2025 - Present**
- Designed Tectonic, a Rust-based workload generator modeling dynamic properties of real-world workloads.
  - Achieved  $2\times$  higher throughput and  $84\%$  lower memory footprint compared to state-of-the-art benchmarks.
- LSM Buffer Analysis** [[Paper PDF](#)] [[Blog](#)] [[GitHub](#)] **Jan 2024 - Present**
- Developed and benchmarked multiple in-memory buffer designs (skiplist, vector, hash-hybrids) for Log Structured Merge Tree (LSM) based storage engines, optimizing throughput, latency, and memory usage.
  - Created a framework for optimal buffer selection for various workloads, enabling data-driven design choices.
- RangeReduce** [[Blog](#)] [[GitHub](#)] **Sep 2023 - Present**
- Designed RangeReduce, a range query-optimized LSM engine that leverages range queries output to derive compactions, improving system performance by reducing disk reads/writes and saving CPU-cycles.
  - Achieved up to  $18\%$  lower range-query latency,  $90\%$  reduction in compaction debt,  $12\%$  less overall data movement, and  $20\%$  lower space utilization, enhancing overall LSM engine efficiency.
- Edge Aware Stream Processing** [[Blog](#)] [[GitHub](#)] **Jan 2023 - May 2023**
- Developed a dynamic operator placement algorithm in Apache Flink to offload tasks to Raspberry Pi edge devices, leveraging performance metrics to reduce WAN data traffic.
  - Prototype preprocessed data streams at the edge, improving system efficiency and demonstrating lower latency and server-side resource usage compared to conventional homogeneous deployments.
- Automated Vulnerability Detection** [[Blog](#)] [[Report](#)] [[GitHub](#)] **Jan 2023 - May 2023**
- Developed an automated tool to detect security risks in third-party VS Code extensions, targeting Path Traversal and Zip Slip attacks, and identified vulnerabilities in  $5\%$  of extensions analyzed.

## SELECTED PUBLICATIONS / POSTERS

DBTest 2024	Shubham Kaushik and Subhadeep Sarkar. <i>Anatomy of the LSM Memory Buffer: Insights &amp; Implications</i> , In Proceedings of the International Workshop on Testing Database Systems
JCSE 2019	Shubham Kaushik and Ratneshwer. <i>Fault Modeling of an Object-Oriented System using CPN</i> , International Journal of Computer Sciences and Engineering
NEDB Day 2025/2024	Shubham Kaushik, Manos Athanassoulis, and Subhadeep Sarkar <i>RangeReduce: A Range Query Driven Compaction for LSM-Trees</i> , North East Database Day (poster)

## CURRICULAR ACTIVITIES

- Member, VLDB 2026 Shadow PC
- Judged and mentored at [HackMIT 2023](#) and [BostonHacks 2022](#)